

PyTorch Fundamentals

Prepared By:

H. Fuat Alsan

PyTorch

- Very Pythonic
- Efficient Tensor operations
- Autograd (Automatic Differentiation)
- Ready-to-use deep learning layers and functions
- Handles complex CUDA operations for us
- Uses cuDNN (CUDA Deep Neural Network)

PyTorch Important Components

- Tensors
 - Represent data and math operations
 - Very similar to NumPy arrays
- Dataset
 - Represent all available data
 - Transforms (data augmentation)
- Dataloader
 - Creates batches from the datasets
- Autograd (Automatic Differentiation)
 - Useful for gradient descent, backpropagation
- Model and Layers (nn.Module)
- Optimizers
- Loss Functions
- Learning Rate (LR) schedulers

Datasets

- `torch.utils.data.Dataset`
- **Built-in Datasets (`torchvision.datasets`)**
 - MNIST
 - CIFAR10
- **Imagefolder**
 - Useful for classification
- **Custom Dataset Class**
 - Most flexible
- See: https://pytorch.org/tutorials/beginner/basics/data_tutorial.html

Datasets

- Train/validation/test split

Data transformation &
augmentations

Download & load built-in
dataset for training and
testing

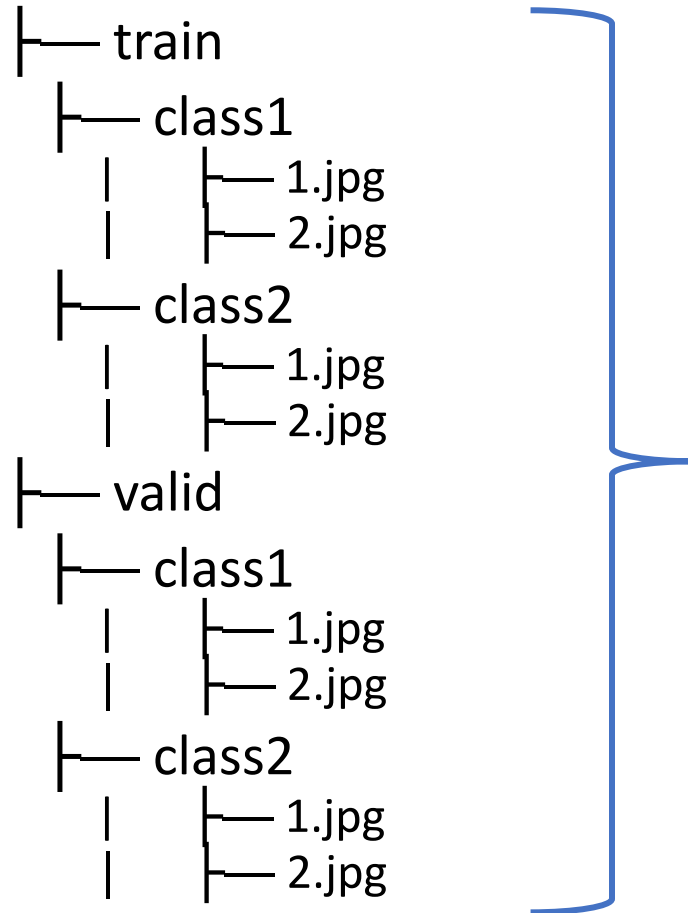
```
from torchvision import datasets, transforms

transform=transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)),
])

train_dataset = datasets.MNIST(
    root='./data',
    train=True,
    download=True,
    transform=transform
)

test_dataset = datasets.MNIST(
    root='./data',
    train=False,
    transform=transform
)
```

ImageFolder



(Root data folder structure)

ImageFolder

root_folder/train/dogs/001.png

root_folder/train/dogs/002.png

...

root_folder/val/dogs/001.png

root_folder/val/dogs/002.png

...

root_folder/train/cats/001.png

root_folder/train/cats/002.png

...

root_folder/val/cats/001.png

root_folder/val/cats/002.png

```
train_dataset = torchvision.datasets.ImageFolder(root='train', train_transforms)
```

```
valid_dataset = torchvision.datasets.ImageFolder(root='val', val_transforms)
```

(AUTOMATICALLY ASSINGS CLASS LABELS USING FOLDER STRUCTURE)

Custom Dataset Class

Subclass of Dataset

Total length of dataset

Function used for
fetching data

Reading image file

Convert to tensors

Apply data augmentation

```
class BasicSegmentationDataset(Dataset):
    def __init__(self, df_full, class_colors, class_mapping, transform=None, reverse=True):
        self.df_full = df_full
        self.class_colors = class_colors
        self.class_mapping = class_mapping
        self.transform = transform
        self.reverse = reverse

        self.xml_files = df_full.filename.unique()

    def __len__(self):
        return len(self.xml_files)

    def __getitem__(self, idx):
        # id_rooms.png (target image)
        selected_xml_file = self.xml_files[idx]
        selected_img_file = selected_xml_file[:-4] + '.tiff'

        image_1 = Image.open(selected_img_file).convert('RGB')
        # for grayscale
        #image_2 = Image.open(os.path.join(self.file_dir, selected_img_file)).convert('L')
        #image_2 = Image.open(os.path.join(self.file_dir, selected_img_file_2)).convert('RGB')

        image_1 = np.array(image_1).astype(np.float32) # ORIGINAL IMAGE
        #image_2 = swap_pixels(image_1.copy()).astype(np.float32) # MASK
        _df = self.df_full[self.df_full.filename == selected_xml_file]
        image_2 = self.draw_masks(image_1.copy(), _df)

        # Feature scaling
        image_1 = normalize_image(image_1)
        image_2 = normalize_image(image_2)

        # convert to tensor
        # (Width, Height, Channel) -> (Channel, Width, Height)
        image_1 = torch.from_numpy(image_1.copy().transpose((2,0,1)))
        image_2 = torch.from_numpy(image_2.copy().transpose((2,0,1)))
        # for grayscale (repeats grayscale channel over RGB channels)
        #image_2 = torch.from_numpy(image_2.copy()).unsqueeze(0).repeat(3, 1, 1)

        # Apply image transformation (if any)
        if self.transform is not None:
            both_images = torch.cat((image_1.unsqueeze(0), image_2.unsqueeze(0)), 0)
            transformed_images = self.transform(both_images)

            image_1 = transformed_images[0]
            image_2 = transformed_images[1]

        return image_1, image_2
```


Dataset Batching

Pass dataset

Batch size

Add batch dimension to the data

Original: (1, 28, 28)

Batched: (32, 1, 28, 28)

```
train_loader = torch.utils.data.DataLoader(  
    dataset=train_dataset,  
    batch_size=32,  
    shuffle=True,  
    num_workers=6,  
    pin_memory=False,  
)
```

```
test_loader = torch.utils.data.DataLoader(  
    dataset=test_dataset,  
    batch_size=32,  
    shuffle=True,  
    num_workers=6,  
    pin_memory=False,  
)
```

```
ex_img_batch, ex_target_batch = next(iter(train_loader))  
print(ex_img_batch.shape)  
print(ex_target_batch.shape)
```

```
torch.Size([32, 1, 28, 28])  
torch.Size([32])
```

PyTorch nn.Module Hierarchy

- **nn.Module** is the base class for everything in PyTorch
- nn.Module2 **(contains)** nn.Module1 **(contains)** nn.Parameter **(contains)** tensors
- Commonly used layers:
 - nn.Linear
 - nn.Conv2d
 - nn.ConvTranspose2d
 - nn.ReLU
 - nn.Sigmoid

Basic nn.Module




Subclass of nn.Module

super().__init__() must be called!

__init__() is special function that defines the module itself

nn.Module can contain other nn.Module(s)

Forward function is called when forward pass on this module is done

```
class SimpleLinear(nn.Module):  
    def __init__(self, in_features, out_features):  
        super().__init__()   
        self.in_features = in_features  
        self.out_features = out_features  
  
        self.fc1 = nn.Linear(in_features, out_features)   
  
        self.act_fn = nn.ReLU()  
  
    def forward(self, x):   
        x = self.fc1(x)  
        x = self.act_fn(x)  
        return x
```

There is no need for defining backward pass function, autograd handles that for us

Another nn.Module

`__init__()` function
takes arguments to
define the module

Other nn.Modules

Forward pass

```
class DownsampleBlock(nn.Module):
    def __init__(self, in_channels, out_channels, down_ratio):
        super().__init__()

        # Conv Layer with kernel_size=(3, 3) and padding=1
        # This Layer normally doesn't change the width and height
        # However, setting stride=down_ratio we downsample the image
        # Also: channels are converted fro in_channels to out_channels
        self.conv = nn.Conv2d(
            in_channels=in_channels,
            out_channels=out_channels,
            kernel_size=3,
            stride=down_ratio, # divide by down_ratio
            padding=1,
            bias=False
        )

        self.norm = nn.BatchNorm2d(out_channels)

        self.relu = nn.ReLU()

    def forward(self, x_in):
        # (N, in_channels, H, W)
        x_features = self.conv(x_in)

        # (N, out_channels, H//down_ratio, W//down_ratio)
        # doesn't change dims, only feature scaling
        x_features = self.norm(x_features)

        # (N, out_channels, H//down_ratio, W//down_ratio)
        # doesn't change dims, only non-linear activation
        x_features = self.relu(x_features)

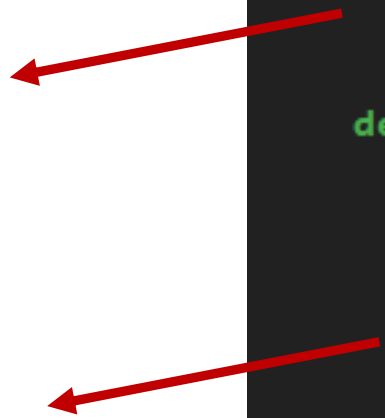
        # final dims: (N, out_channels, H//down_ratio, W//down_ratio)
        return x_features
```

nn.Module with No Learnable Parameters

Does not have any
other nn.Module(s)

Performs a very
basic forward pass
operation

```
class Flatten(nn.Module):  
    def __init__(self):  
        super().__init__()  
  
    def forward(self, x):  
        # (N, CHANNELS, HEIGHT, WIDTH)  
        batch_size = x.shape[0]  
        # Merge into single dimension  
        # (N, CHANNELS*HEIGHT*WIDTH)  
        x = x.reshape(batch_size, -1)  
        return x
```



nn.Module with No Learnable Parameters

```
class UpscaleLayer(nn.Module):
    def __init__(self, scale_factor):
        super().__init__()
        self.scale_factor = scale_factor

    def forward(self, x):
        # (N, CHANNELS, HEIGHT, WIDTH)
        # Perform nearest-neighbor interpolation to upscale the input
        # Assigns the value of the nearest input pixel
        x_upscaled = nn.functional.interpolate(x, scale_factor=self.scale_factor, mode='nearest')
        # x_upscaled dims:
        # (N, CHANNELS, HEIGHT*scale_factor, WIDTH*scale_factor)
        return x_upscaled
```

Same as before but different forward pass operation

Full Model

Each of these are
nn.Module layers

Convolutional part

Fully connected
(nn.Linear) part

```
class ConvNet(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(in_channels=1, out_channels=32, kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3)

        self.max_pool = nn.MaxPool2d(kernel_size=2)

        self.relu = nn.ReLU()

        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        #####
        # Convolutional Part
        #####
        print(f'Input dims: {x.shape}')
        x = self.conv1(x) # (N, 1, 28, 28) -> (N, 32, 26, 26)
        print(f'After conv1 {x.shape}')
        x = self.relu(x) # no dim change
        x = self.conv2(x) # (N, 32, 26, 26) -> (N, 64, 24, 24)
        print(f'After conv2 {x.shape}')
        x = self.relu(x) # no dim change
        x = self.max_pool(x) # (N, 64, 24, 24) -> (N, 64, 12, 12)
        print(f'After maxpool {x.shape}')
        #####
        #####

        #####
        ## Fully Connected Part
        #####
        x = torch.flatten(x, 1) # (N, 64, 12, 12) -> (N, 64*12*12) -> (N, 9216)
        x = self.fc1(x) # (N, 9216) -> (N, 128)
        x = self.relu(x) # no dim change
        logits = self.fc2(x) # (N, 128) - (N, 10)
        #####
        #####

        return logits
```

Classification Model

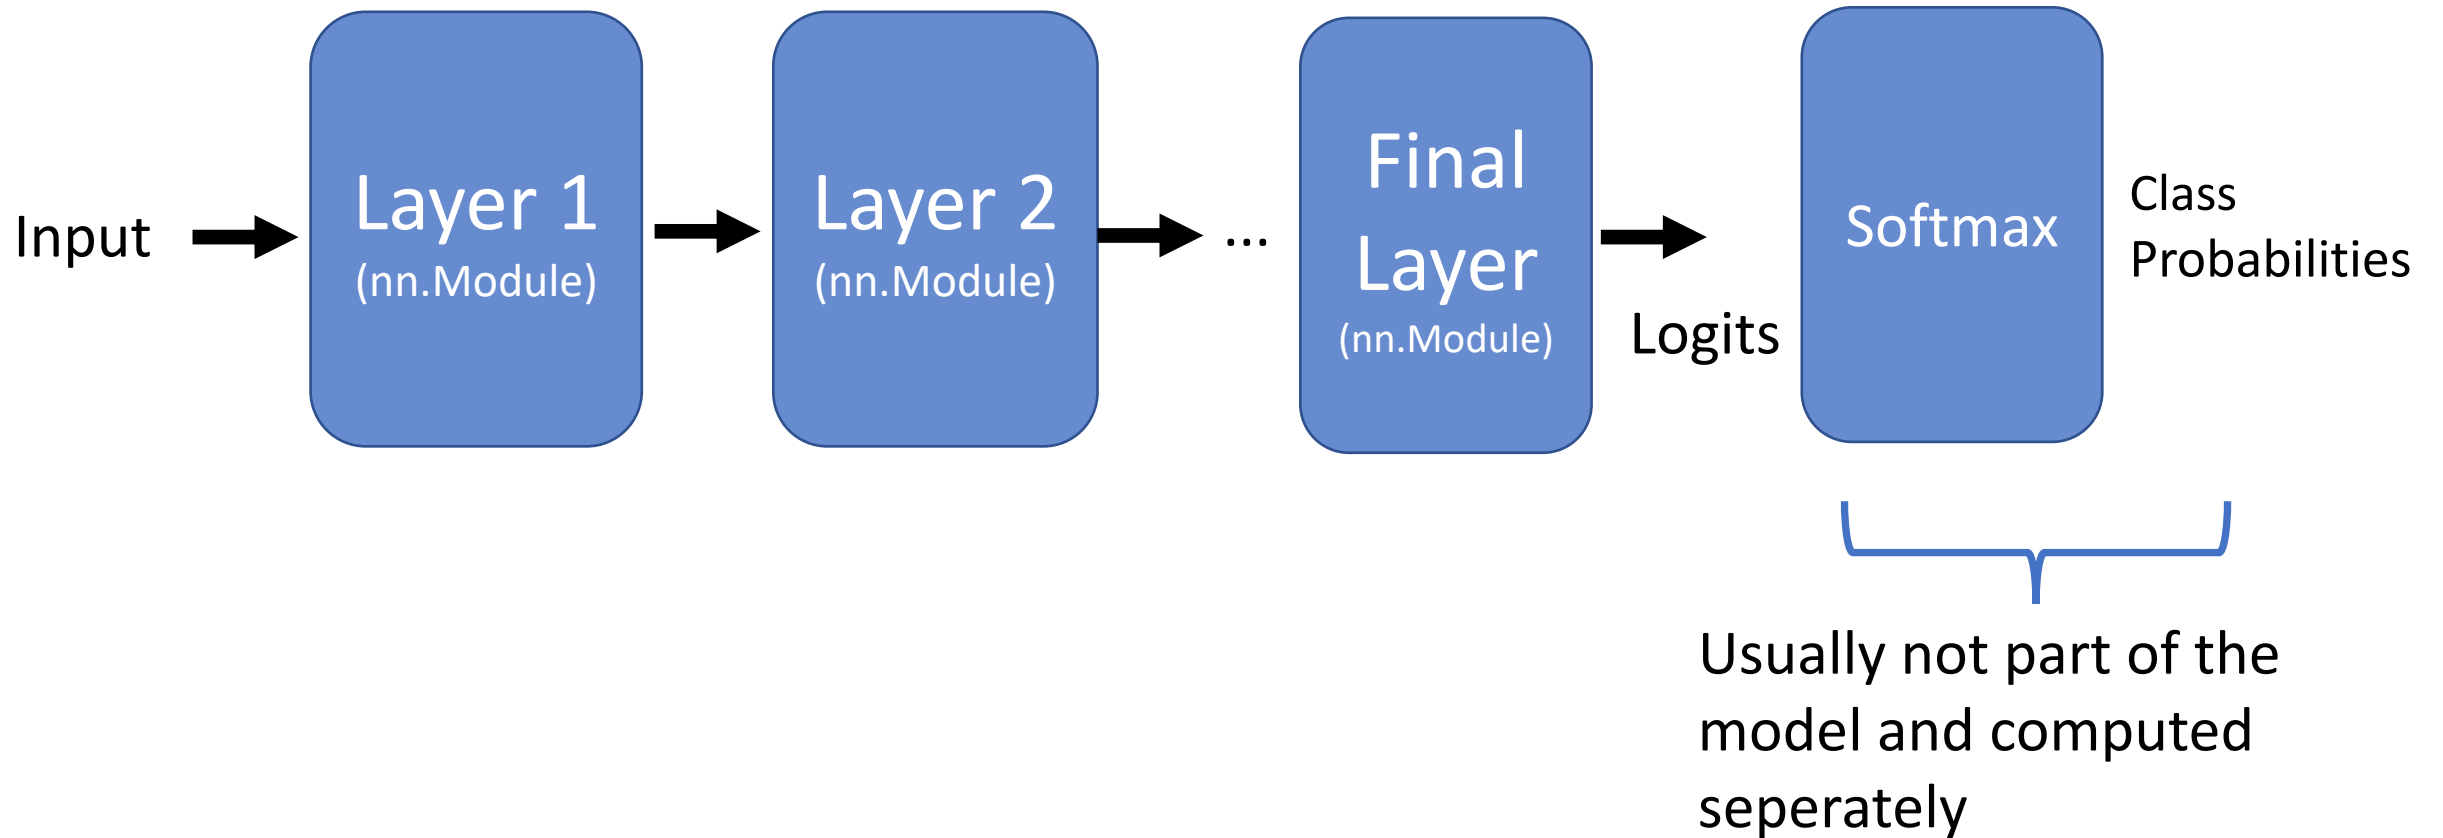
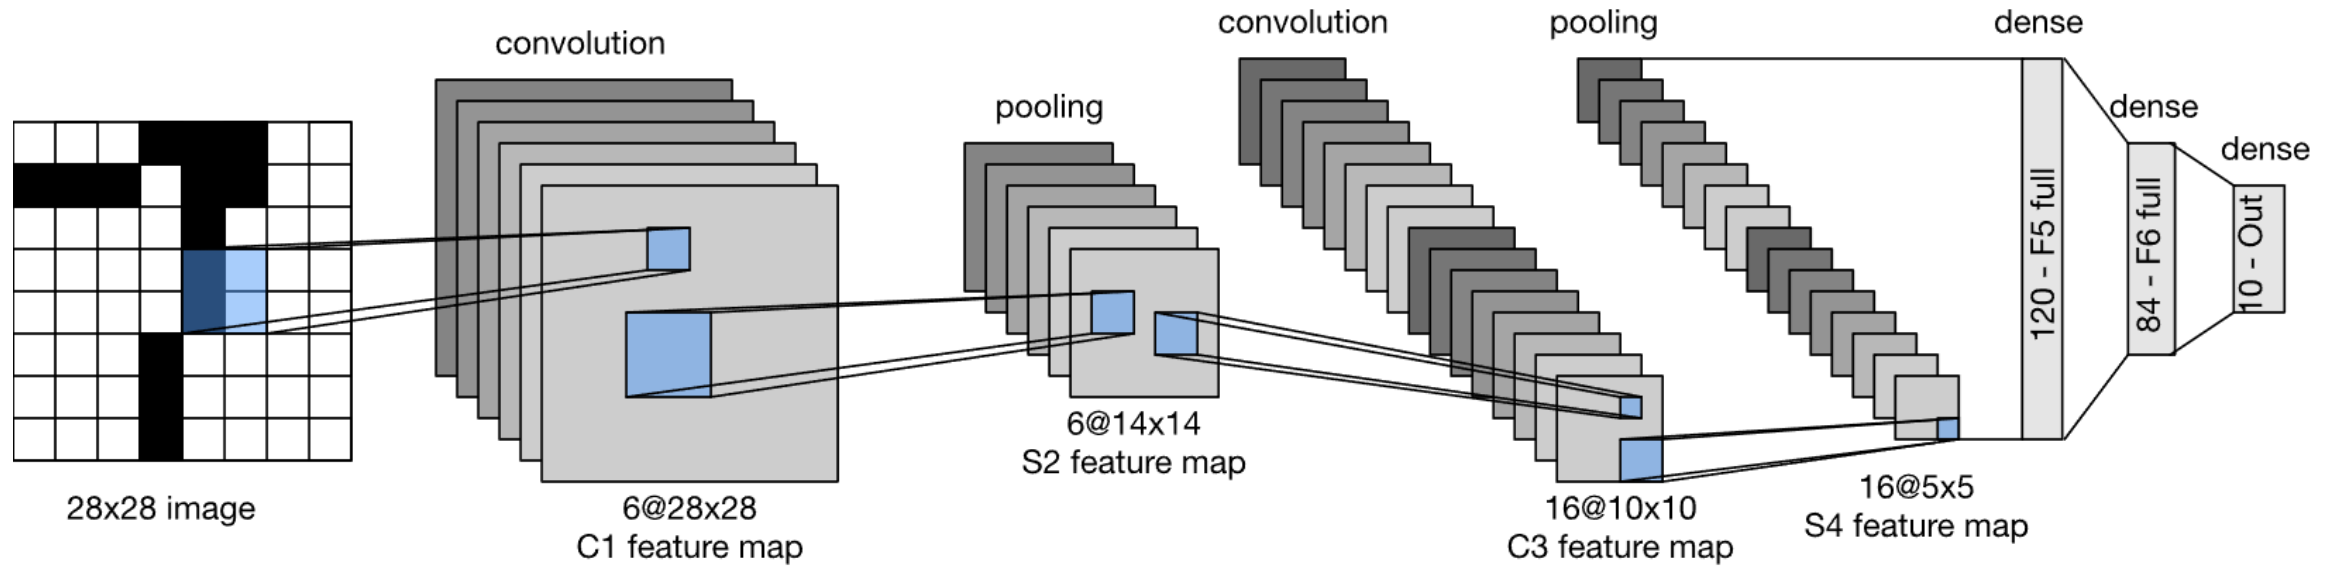
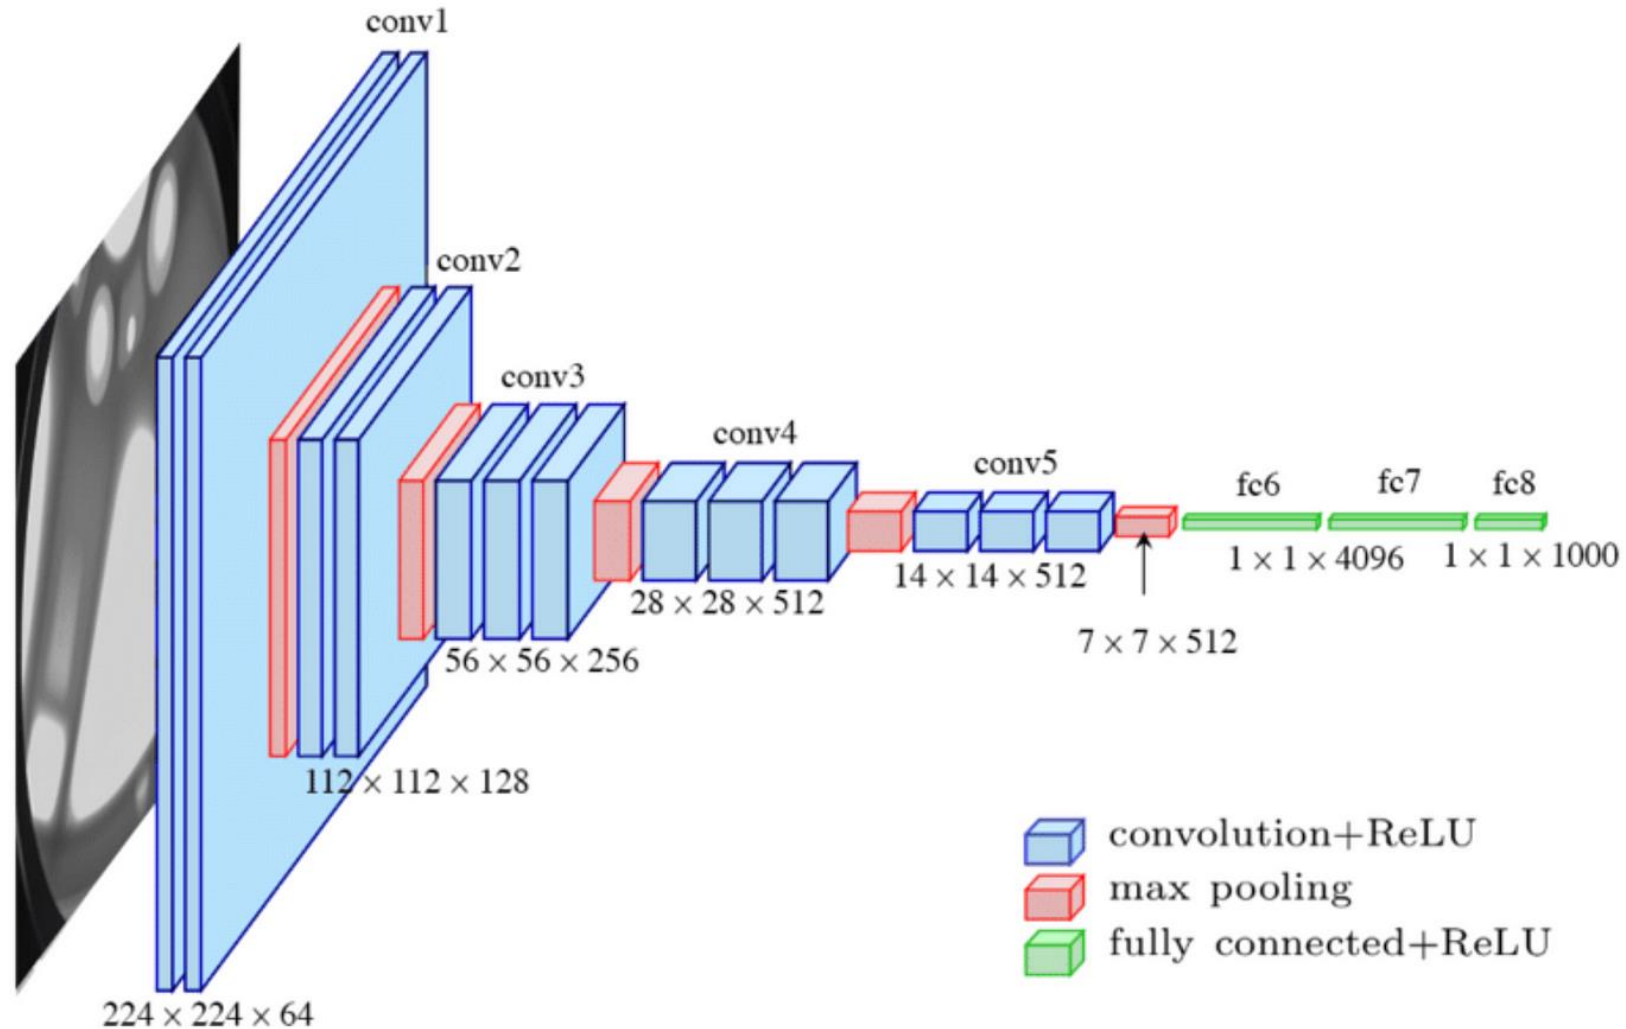


Image Classifier for MNIST Dataset



A Modern Image Classifier (VGG16)



Training Classifiers

- In regression we had continuous values, we used MSE
- However, this is not the case in classification
- In classification, we have fixed classes
 - Each class is represented by an integer ID
 - For ex: Cat/Dog classifier
 - Cat \rightarrow 0
 - Dog \rightarrow 1
- In classification, we use a loss function called **cross entropy**
 - Actually it is **negative log-likelihood** loss

Training Classifiers

- **Batch iteration:** single cycle of batch data
 - **batch_size** amount of data is used for every iteration
 - Last iteration might have less data (dataset is not fully divisible by batch size)
 - (Use **drop_last=True** in dataloader to ignore this)
- **Epoch:** single cycle of full dataset
 - All batch iterations are completed
- **Backpropagation**
 - **Foward pass:** Model predicts, loss is computed
 - **Backward pass:** Using loss, gradients are computed and parameters are updated

Training Loop

Switch model to
training mode

Loop over all batches

Move data to GPU
(if available)

Forward pass and
compute loss

Backward pass,
compute gradients
and update model
parameters

```
def train(model, train_loader, optimizer, criterion, epoch):  
    model.train()  
  
    for batch_idx, (img, target) in enumerate(train_loader):  
        img, target = img.to(device), target.to(device)  
  
        # Zero gradients, perform a backward pass, and update the weights.  
        # In PyTorch, gradients are accumulated, you need to reset gradients in each loop  
        optimizer.zero_grad()  
  
        # Forward pass  
        preds = model(img)  
        loss = criterion(preds, target)  
  
        # Compute gradients  
        loss.backward()  
        # Update gradients  
        optimizer.step()
```

Test/Validation Loop

Disable Autograd

Switch model to eval mode

Loop over all batches

Move data to GPU
(if available)

Forward pass and
compute loss

Compute accuracy

There is no backward pass
in testing/validation

```
@torch.no_grad()
def test(model, test_loader, criterion):
    model.eval()

    test_loss = 0
    correct = 0

    for img, target in test_loader:
        img, target = img.to(device), target.to(device)

        preds = model(img)
        test_loss += criterion(preds, target)

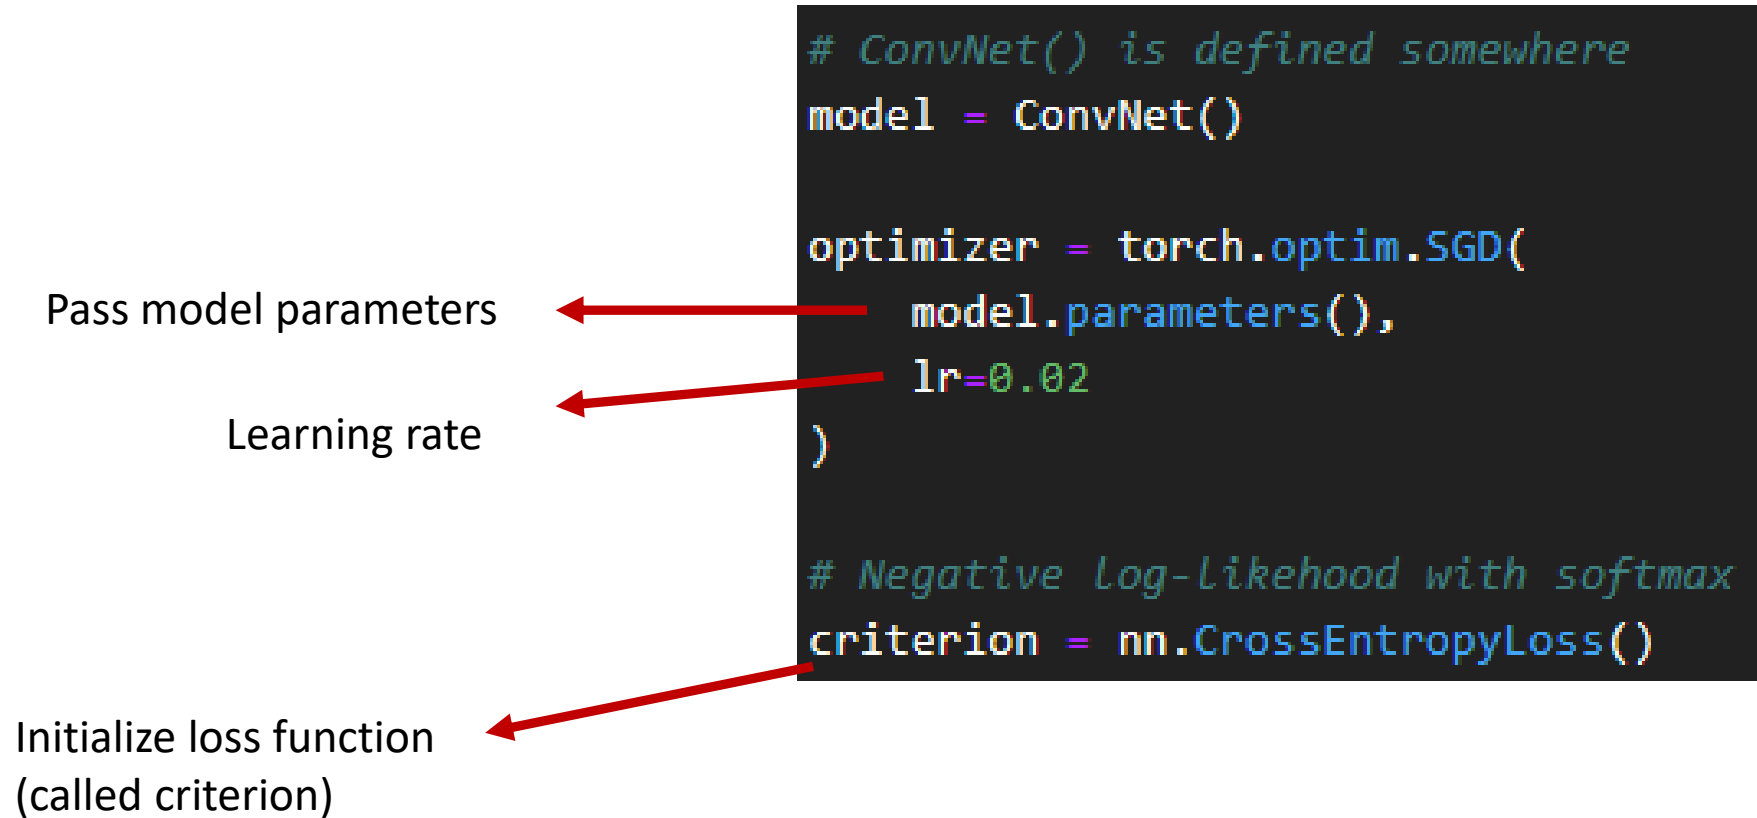
        pred_max = preds.argmax(dim=1, keepdim=True) # get the index of the max probable class
        correct += pred_max.eq(target.view_as(pred_max)).sum().item()

    test_loss /= len(test_loader.dataset)

    test_acc = 100.0 * correct / len(test_loader.dataset)

    return test_loss
```

Loss & Optimizer



Epoch Loop

```
NUM_EPOCHS = 10

# Move model to GPU (if available)
model.to(device)

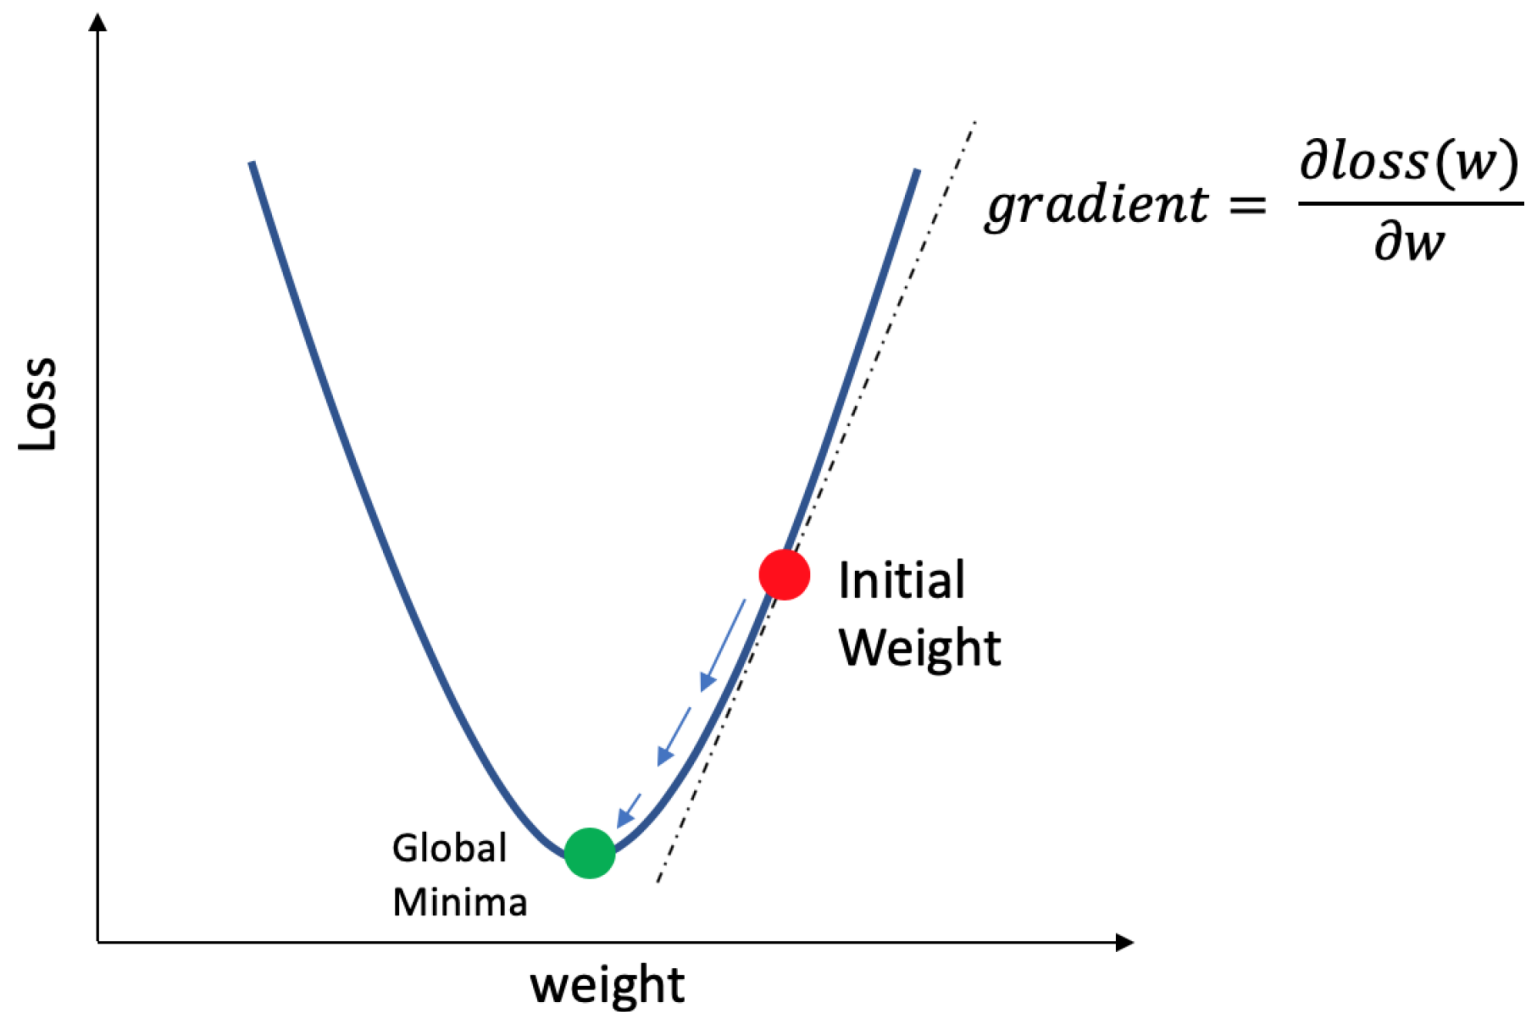
for epoch in range(1, NUM_EPOCHS+1):
    train_loss, train_acc = train(model, train_loader, optimizer, criterion, epoch)
    test_loss, test_acc = test(model, test_loader, criterion)

    print(f'Epoch: {epoch}, Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.2f}, Test Loss: {test_loss:.4f}, Test Acc: {test_acc:.2f}')
```

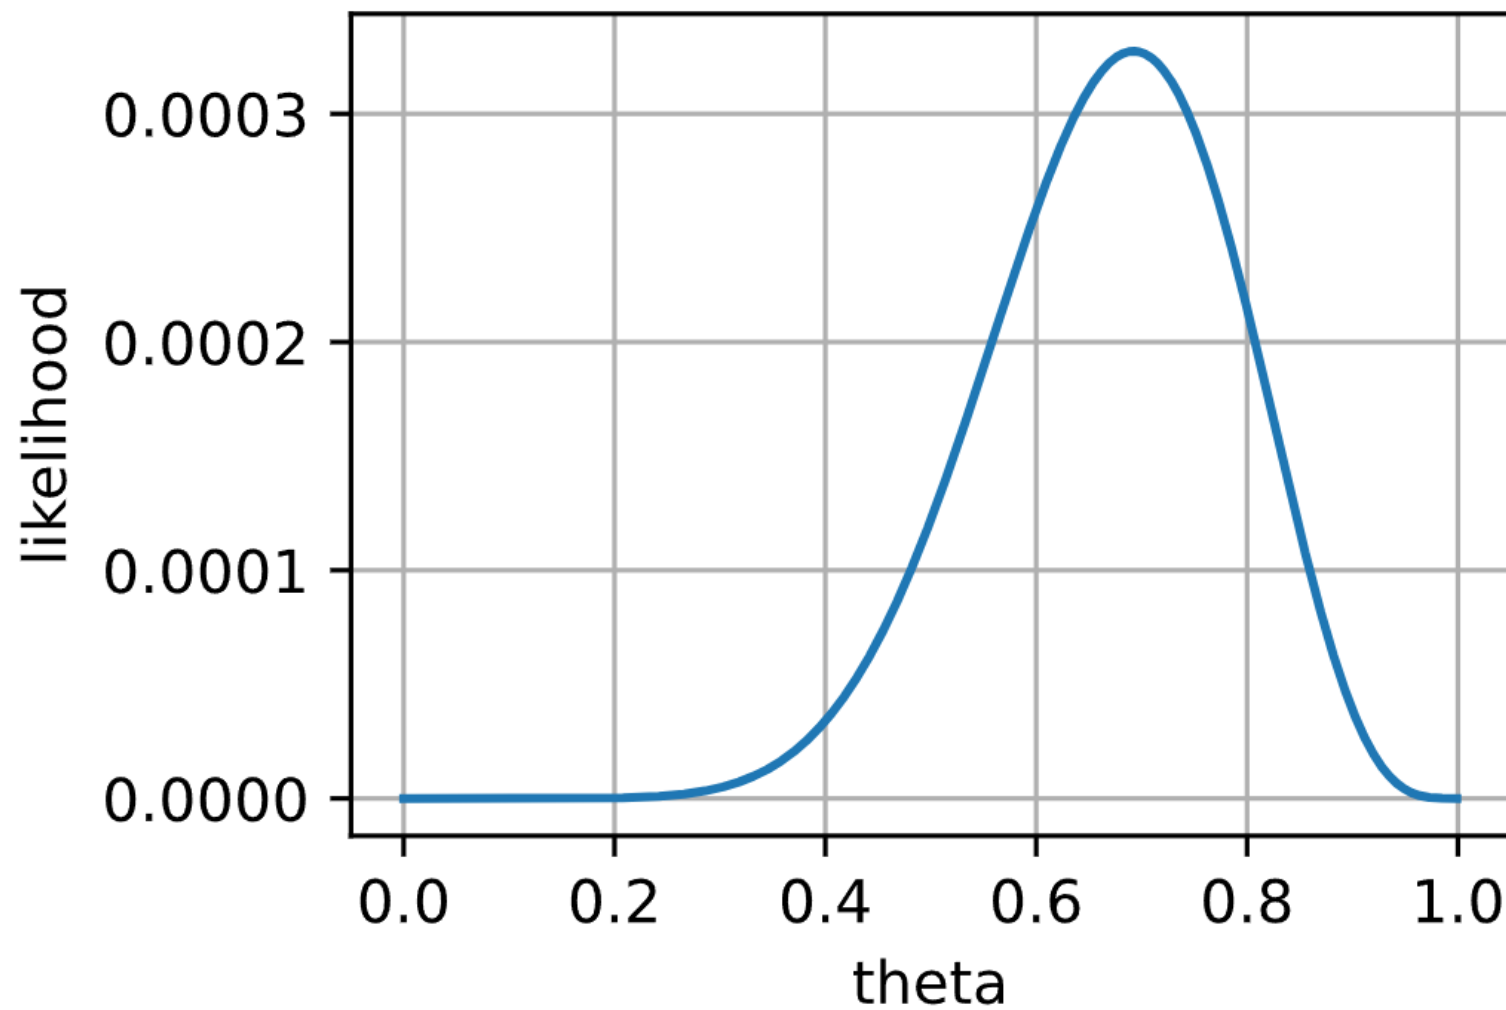

Likelihood Function As Loss

- We wish to increase the probability of predicting a class
- Likelihood function is useful for measuring that
- Log likelihood: turns multiplication into addition
 - Faster computation -> Negative log likelihood
- Negative log likelihood
 - Likelihood: better when increased
 - Gradient descent is designed for reducing
 - Hence we add «negative sign» to log likelihood
 - This way, we increase the likelihood

Loss vs Likelihood



Loss vs Likelihood



Classification Loss Function

- Binay Case:
- y_i : class id (0 or 1, each indicating a class)
- \hat{y}_i : probability of class (predicted by the model)

$$\text{Loss} = -\frac{1}{\text{output size}} \sum_{i=1}^{\text{output size}} y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log (1 - \hat{y}_i)$$

BCE PyTorch Implementation

- Binary Cross Entropy (BCE)
- Automatically flattens the data if `tensor.dim() > 2`
 - For ex: segmentation mask have (N, 10, height, width)
 - Useful for computing loss on segmentation masks
 - (N, flattened_mask_actual), (N, flattened_mask_predected)

Multiclass Classification (Cont.)

$$\text{logloss} = -\frac{1}{N} \sum_i^N \sum_j^M y_{ij} \log(p_{ij})$$

- N is the number of rows
- M is the number of classes

Save & Load Model

Model parameters represented as
Python dictionary

Optimizer inner state represented as
Python dictionary

Filename

Load model state dict

Do this only if training will NOT continue

```
# SAVE MODEL AND OPTIMIZER STATE DICT
# SAVE FILENAME 'convnet_checkpoint.pt'
torch.save({
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict()
},
    'convnet_checkpoint.pt'
)
```

```
# LOAD PRE-TRAINED MODEL
model = ConvNet()

checkpoint = torch.load('convnet_checkpoint.pt')
model.load_state_dict(checkpoint['model_state_dict'], strict=True)

# SWITCH MODEL TO PREDICTION ONLY MODE
# (OPTIONAL)
model.eval()
```